

K-Means Clustering

what is K-Means?

- given $K > 0$ (where K is the number of clusters)
- a set of N d -dimensional objects
 - Clustering is the process of grouping a set of N d -dimensional (2-dimensional, 3-dimensional, etc.) objects into K clusters of similar objects.
 - Objects should be similar to one another within the same cluster and dissimilar to those in other clusters.
- K-Means is a distance-based clustering algorithm.
- K-Means clustering (aka unsupervised learning) is a data mining algorithm to cluster, classify, or group your N objects based on their attributes or features into K number of groups (so-called clusters).

For example

- It can be used to find a group of consumers with common behaviors.
- To cluster documents based on the similarity of their contents
- Selection of K is specific to the application or problem domain.

K-Means algorithm

1. Partition N objects into K nonempty subsets.
2. Compute the centroids of the clusters in the current partition (the centroid is the center, or mean point, of the cluster). For all $i = 1, 2, \dots, K$, compute μ_i as:

$$\mu_i = \frac{1}{|c_i|} \sum_{j \in c_i} x_j \quad \forall i$$

3. Assign each object to the cluster with the nearest centroid. This is simply attributing the closest cluster to each data point:

$$c_i = \{j : d(x_j, \mu_i) \leq d(x_j, \mu_l), l \neq i, j = 1, \dots, n\}$$

where $d(a, b)$ is the distance function for two points: a and b .

4. Stop when there are no more new assignments. Otherwise, go back to step 2. Basically, we repeat steps 2 and 3 until convergence.

K-Means Distance Function

- Euclidean distance, Manhattan distance, Inner product space, Maximum norm, Your own custom function (any metric you define over the d -dimensional space)

Let:

$$X = (X_1, X_2, \dots, X_d)$$

$$Y = (Y_1, Y_2, \dots, Y_d)$$

then:

$$\text{distance}(X, Y) = \sqrt{(X_1 - Y_1)^2 + (X_2 - Y_2)^2 + \dots + (X_d - Y_d)^2}$$

The Euclidean distance function has some interesting properties:

- $\text{distance}(i, j) \geq 0$
- $\text{distance}(i, i) = 0$
- $\text{distance}(i, j) = \text{distance}(j, i)$
- $\text{distance}(i, j) \leq \text{distance}(i, k) + \text{distance}(k, j)$

MapReduce Solution for K-Means Clustering

Example 12-1. K-Means clustering algorithm

```
1 // k = number of desired clusters
2 // delta = acceptable error for convergence
3 // data = input data
4 kmeans(k, delta, data) {
5     // initialize the cluster centroids
6     initial_centroids = pick(k, data);
7
8     // this is how we broadcast centers to mappers
9     writeToHDFS(initial_centroids);
10
11     // iterate as long as necessary
12     current_centroids = initial_centroids;
13     while (true) {
14         // theMapReduceJob() does 2 tasks:
15         //     1. uses current_centroids in map()
16         //     2. reduce() creates new_centroids and writes it to HDFS
17         theMapReduceJob();
18         new_centroids = readFromHDFS();
19         if change(new_centroids, current_centroids) <= delta {
20             // we are done, terminate loop-iteration
21             break;
22         }
23         else {
24             current_centroids = new_centroids;
25         }
26     }
27
28     result = readFromHDFS();
29     return result;
30 }
```

Example 12-2. K-Means clustering algorithm: change() method

```
1 change(new_centroids, current_centroids) {
2     new_distance = [sum of squared distance in the new_centroids];
3     current_distance = [sum of squared distance in the current_centroids];
4     changed = absoluteValue(new_distance - current_distance);
5     return changed;
6 }
```

MapReduce Solution: map()

- Read the cluster centroids into memory from a SequenceFile1 (note that we may also use Redis or memcached for persisting cluster centroids). In Hadoop's implementation, this will be done in the setup() method of the mapper class.
- Iterate over each cluster centroid for each input key-value pair. In Hadoop's implementation, for the map() function, the key is generated by Hadoop and ignored (not used).
- Compute the Euclidean distances and save the nearest center with the lowest distance to the input point (as a d -dimensional vector).
- Write the key-value pair to be consumed by reducers, where the key is the nearest cluster center to the input point (and the value is a d -dimensional vector). Both the key and the value are of the Vector data type.

MapReduce Solution: map()

Example 12-3. K-Means MapReduce: map() function

```
1 public class KmeansMapper ... {
2
3     private List<Vector> centers = null;
4
5     private List<Vector> readCentersFromSequenceFile() {
6         // read cluster centroids from a SequenceFile,
7         // which is a set of key-value pairs
8         ...
9     }
10
11     // called once at the beginning of the map task
12     public void setup(Context context) {
13         this.centers = readCentersFromSequenceFile();
14     }
15
16     * @param key is MapReduce generated, ignored here
17     * @param value is the d-dimensional Vector (V1, V2, ..., Vd)
18     */
19
20     map(Object key, Vector value) {
21         Vector nearest = null;
22         double nearestDistance = Double.MAX_VALUE;
23         for (Vector center : centers) {
24             double distance = EuclideanDistance.calculateDistance(center, value);
25             if (nearest == null) {
26                 nearest = center;
27                 nearestDistance = distance;
28             }
29             else {
30                 if (nearestDistance > distance) {
31                     nearest = center;
32                     nearestDistance = distance;
33                 }
34             }
35         }
36
37         // prepare key-value for reducers
38         emit(nearest, value);
39     }
}
```


MapReduce Solution: combine()

```
1 /**
2  * @param key is the Centroid
3  * @param values is a list of Vectors
4  */
5 combine(Vector key, Iterable<Vector> values) {
6     // all dimensions in sum Vector are initialized to 0.0
7     Vector sum = new Vector();
8     for (Vector value : values) {
9         // note that value.length = d,
10        // where d is the number of dimensions for input objects
11        for (int i = 0; i < value.length; i++) {
12            sum[i] += value[i];
13        }
14    }
15
16    emit(key, sum);
17 }
```

MapReduce Solution: reduce()

- The main task of the reduce() function is to recenter.
- Each reducer iterates over each value vector and calculates the mean vector. Once you have found the mean, this is the new center; your final step is to save it into a persistent store (such as a SequenceFile).

```
1 /**
2  * @param key is the Centroid
3  * @param values is a list of Vectors
4  */
5 reduce(Vector key, Iterable<Vector> values) {
6     // all dimensions in newCenter are initialized to 0.0
7     Vector newCenter = new Vector();
8     int count = 0;
9     for (Vector value : values) {
10         count++;
11         for (int i = 0; i < value.length; i++) {
12             newCenter[i] += value[i];
13         }
14     }
15     for (int i = 0; i < key.length; i++) {
16         // set new mean for each dimension
17         newCenter[i] = newCenter[i] / count;
18     }
19     emit(key.ID, newCenter);
20 }
21
22 }
```